

predictive_project

2024-03-18

Predictive Project

Read data

```
vehicles <- read.csv("~/Downloads/vehicles.csv")
tail(vehicles)
```

```
##           id
## 426875 7301591199
## 426876 7301591192
## 426877 7301591187
## 426878 7301591147
## 426879 7301591140
## 426880 7301591129
##
##                                     url
## 426875 https://wyoming.craigslist.org/ctd/d/atlanta-2018-lexus-gs-gs-350-sedan-4d/7301591199.html
## 426876 https://wyoming.craigslist.org/ctd/d/atlanta-2019-nissan-maxima-sedan-4d/7301591192.html
## 426877 https://wyoming.craigslist.org/ctd/d/atlanta-2020-volvo-s60-t5-momentum/7301591187.html
## 426878 https://wyoming.craigslist.org/ctd/d/atlanta-2020-caddy-cadillac-xt4-sport/7301591147.html
## 426879 https://wyoming.craigslist.org/ctd/d/atlanta-2018-lexus-es-es-350-sedan-4d/7301591140.html
## 426880 https://wyoming.craigslist.org/ctd/d/atlanta-2019-bmw-series-430i-gran-coupe/7301591129.html
##
##           region            region_url price year manufacturer
## 426875 wyoming https://wyoming.craigslist.org 33590 2018      lexus
## 426876 wyoming https://wyoming.craigslist.org 23590 2019      nissan
## 426877 wyoming https://wyoming.craigslist.org 30590 2020      volvo
## 426878 wyoming https://wyoming.craigslist.org 34990 2020      cadillac
## 426879 wyoming https://wyoming.craigslist.org 28990 2018      lexus
## 426880 wyoming https://wyoming.craigslist.org 30590 2019      bmw
##
##           model condition cylinders fuel odometer
## 426875      gs 350 sedan 4d      good 6 cylinders    gas   30814
## 426876      maxima s sedan 4d      good 6 cylinders    gas   32226
## 426877 s60 t5 momentum sedan 4d      good              gas   12029
## 426878      xt4 sport suv 4d      good              diesel  4174
## 426879      es 350 sedan 4d      good 6 cylinders    gas   30112
## 426880 4 series 430i gran coupe      good              gas   22716
##
##           title_status transmission          VIN drive size      type
## 426875      clean      automatic JTHBZ1BLXJA012999   rwd      sedan
## 426876      clean              other 1N4AA6AV6KC367801   fwd      sedan
## 426877      clean              other 7JR102FKXLG042696   fwd      sedan
## 426878      clean              other 1GYFZFR46LF088296              hatchback
## 426879      clean              other 58ABK1GG4JU103853   fwd      sedan
## 426880      clean              other WBA4J1C58KBM14708   rwd      coupe
```

```
##           paint_color
## 426875         white
## 426876
## 426877         red
## 426878         white
## 426879         silver
## 426880
##                                     image_url
## 426875 https://images.craigslist.org/00I0I_hJHfjCUppaEz_0gw0co_600x450.jpg
## 426876 https://images.craigslist.org/00o0o_iiraFnHg8qUz_0gw0co_600x450.jpg
## 426877 https://images.craigslist.org/00x0x_15sbgnxCISvz_0gw0co_600x450.jpg
## 426878 https://images.craigslist.org/00L0L_farM7bxnxRiz_0gw0co_600x450.jpg
## 426879 https://images.craigslist.org/00z0z_bKnIVGLkDTcz_0gw0co_600x450.jpg
## 426880 https://images.craigslist.org/00Y0Y_1EUocjyRxaJz_0gw0co_600x450.jpg
##
## 426875                                     Carvana is the s
## 426876
## 426877                                     Carvana is the safer way to
## 426878
## 426879 Carvana is the safer way to buy a car During these uncertain times, Carvana is dedicated to e
## 426880
##           county state      lat      long      posting_date
## 426875      NA     wy 33.77921 -84.41181 2021-04-04T03:21:34-0600
## 426876      NA     wy 33.78650 -84.44540 2021-04-04T03:21:31-0600
## 426877      NA     wy 33.78650 -84.44540 2021-04-04T03:21:29-0600
## 426878      NA     wy 33.77921 -84.41181 2021-04-04T03:21:17-0600
## 426879      NA     wy 33.78650 -84.44540 2021-04-04T03:21:11-0600
## 426880      NA     wy 33.77921 -84.41181 2021-04-04T03:21:07-0600
```

```
# check number of rows and columns
rows <- nrow(vehicles)
col <- ncol(vehicles)
print(paste("rows",rows))
```

```
## [1] "rows 426880"
```

```
print(paste("col",col))
```

```
## [1] "col 26"
```

step.1 : Check for null or empty cells in dataset

```
# Checking the null values of each feature by plotting
library(ggplot2)

feature_miss_count <- colSums(is.na(vehicles) | vehicles == "")
print(data.frame(feature_miss_count))
```

```
##           feature_miss_count
```

```
## id 0
## url 0
## region 0
## region_url 0
## price 0
## year 1205
## manufacturer 17646
## model 5276
## condition 174104
## cylinders 177678
## fuel 3013
## odometer 4400
## title_status 8242
## transmission 2556
## VIN 161042
## drive 130567
## size 306361
## type 92858
## paint_color 130203
## image_url 68
## description 69
## county 426880
## state 0
## lat 6549
## long 6549
## posting_date 68
```

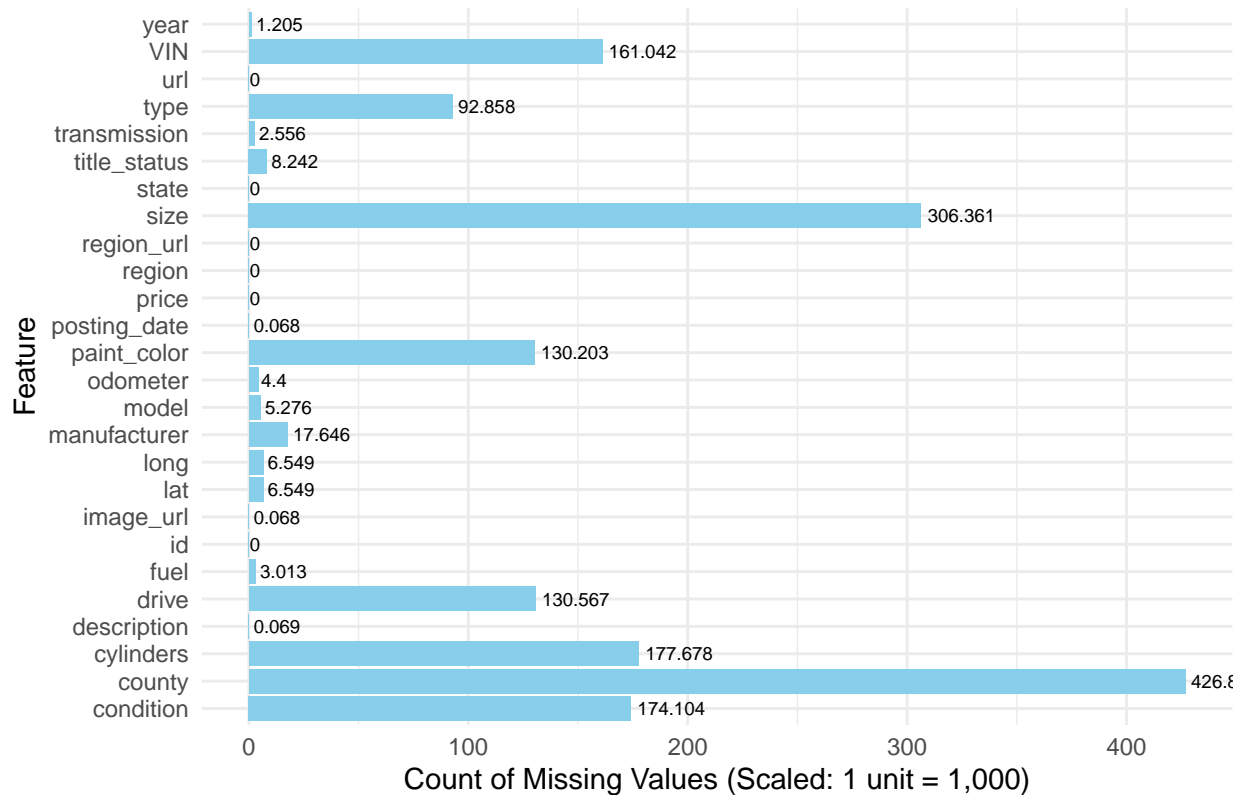
```
scaling <- 1000 # 1 unit represents 1,000 missing values
scaled_miss <- feature_miss_count / scaling
```

```
missing_df <- data.frame(
  Feature = names(scaled_miss),
  Scaled_Missing_Count = scaled_miss
)
```

```
missing_df <- missing_df[order(-missing_df$Scaled_Missing_Count), ]
```

```
ggplot(missing_df, aes(x = Scaled_Missing_Count, y = Feature)) +
  geom_bar(stat = "identity", fill = "skyblue") +
  geom_text(aes(label = Scaled_Missing_Count), vjust = 0.5, color = "black", size=2.6, hjust = - 0.1) +
  labs(
    title = "Count of Missing Values by Feature",
    x = "Count of Missing Values (Scaled: 1 unit = 1,000)",
    y = "Feature"
  ) +
  theme_minimal()
```

Count of Missing Values by Feature



Step2 : Dropping unwanted columns

*# We can see that few columns are unwanted for our use case. So we are dropping them
columns are 'county', 'url', 'region_url', 'VIN', 'image_url', 'region', 'description', 'model'*

```
vehicles <- vehicles[, -c(2,3,4,8,15,20,21,22)]
```

Step3 : Filling missing values

*# If there is any feature with numeric as its datatype, we are replacing with mean of that feature
Else if there is any feature with character as its datatype, we are replacing with mode of that feature*

```
for (i in names(vehicles)[!names(vehicles) %in% c('year', 'manufacturer', 'paint_color')]) {
  if (class(vehicles[[i]]) == "numeric" || class(vehicles[[i]]) == "integer") {
    non_empty_values <- as.numeric(vehicles[[i]][vehicles[[i]] != ""])
    mean_val <- mean(non_empty_values, na.rm = TRUE)
    vehicles[[i]][is.na(vehicles[[i]]) | vehicles[[i]] == ""] <- mean_val
  }
  if (class(vehicles[[i]]) == "character") {
    non_empty_values <- vehicles[[i]][vehicles[[i]] != ""]
    mode_val <- names(sort(table(non_empty_values), decreasing = TRUE))[1] # Find the mode
    vehicles[[i]][is.na(vehicles[[i]]) | vehicles[[i]] == ""] <- mode_val
  }
}
```

```
# Handling missing values for specific columns (i.e year,manufacturer,paint_color)
vehicles$year[is.na(vehicles$year) | vehicles$year == ""] <- names(sort(table(vehicles$year), decreasing=T))
vehicles$manufacturer[is.na(vehicles$manufacturer) | vehicles$manufacturer == ""] <- "Unknown"
vehicles$paint_color[is.na(vehicles$paint_color) | vehicles$paint_color == ""] <- "Unknown"
```

```
head(vehicles)
```

```
##           id price year manufacturer condition  cylinders fuel odometer
## 1 7222695916  6000 2017      Unknown      good 6 cylinders  gas 98043.33
## 2 7218891961 11900 2017      Unknown      good 6 cylinders  gas 98043.33
## 3 7221797935 21000 2017      Unknown      good 6 cylinders  gas 98043.33
## 4 7222270760  1500 2017      Unknown      good 6 cylinders  gas 98043.33
## 5 7210384030  4900 2017      Unknown      good 6 cylinders  gas 98043.33
## 6 7222379453  1600 2017      Unknown      good 6 cylinders  gas 98043.33
## title_status transmission drive      size  type paint_color state      lat
## 1      clean      automatic 4wd full-size sedan      Unknown      az 38.49394
## 2      clean      automatic 4wd full-size sedan      Unknown      ar 38.49394
## 3      clean      automatic 4wd full-size sedan      Unknown      fl 38.49394
## 4      clean      automatic 4wd full-size sedan      Unknown      ma 38.49394
## 5      clean      automatic 4wd full-size sedan      Unknown      nc 38.49394
## 6      clean      automatic 4wd full-size sedan      Unknown      ny 38.49394
##           long           posting_date
## 1 -94.7486 2021-04-23T22:13:05-0400
## 2 -94.7486 2021-04-23T22:13:05-0400
## 3 -94.7486 2021-04-23T22:13:05-0400
## 4 -94.7486 2021-04-23T22:13:05-0400
## 5 -94.7486 2021-04-23T22:13:05-0400
## 6 -94.7486 2021-04-23T22:13:05-0400
```

```
# Box plots before and after cleaning the data
plot_boxplots <- function(data) {
  numeric_cols <- sapply(data, is.numeric)
  par(mfrow = c(3, 3), plt = c(0.05, 1, 0.05, 1)) # Adjust rows, columns, and margins as needed
  for (col in names(data)[numeric_cols]) {
    boxplot(data[[col]], main = col)
    legend("topright", legend = col, bty = "n")
  }
}
```

```
# Plot boxplots for all numeric features
plot_boxplots(vehicles)
```

```
# Function to remove outliers across all features
remove_outliers_all <- function(data) {
  for (col in names(data)) {
    if (is.numeric(data[[col]])) {
      q1 <- quantile(data[[col]], 0.25)
      q3 <- quantile(data[[col]], 0.75)
      iqr <- q3 - q1
      lower_bound <- q1 - 1.5 * iqr
      upper_bound <- q3 + 1.5 * iqr
```

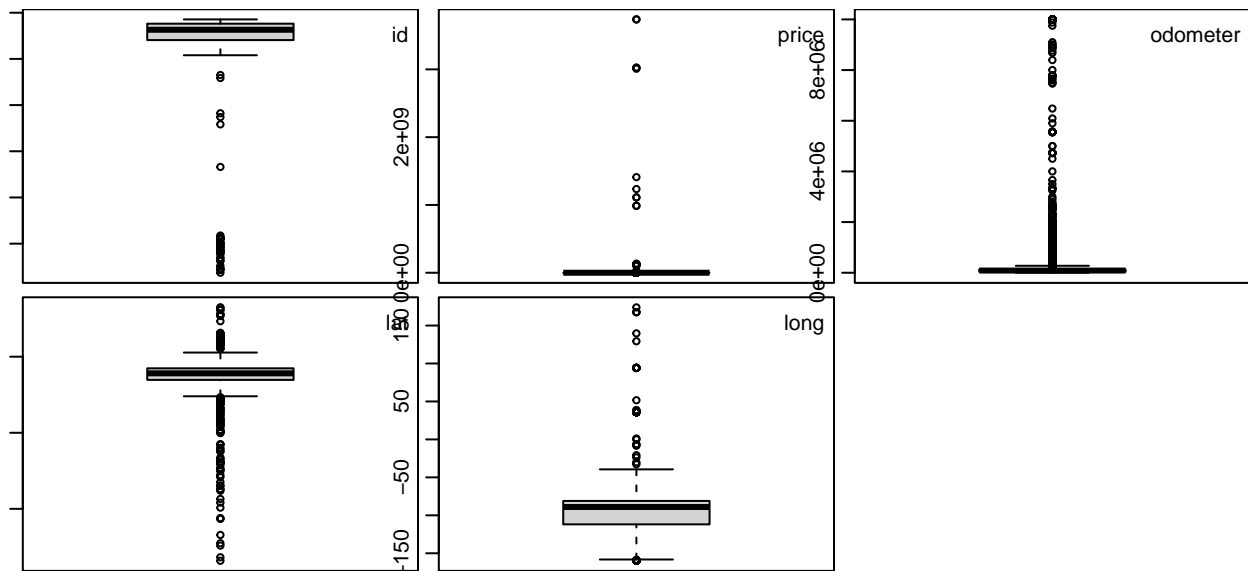
```

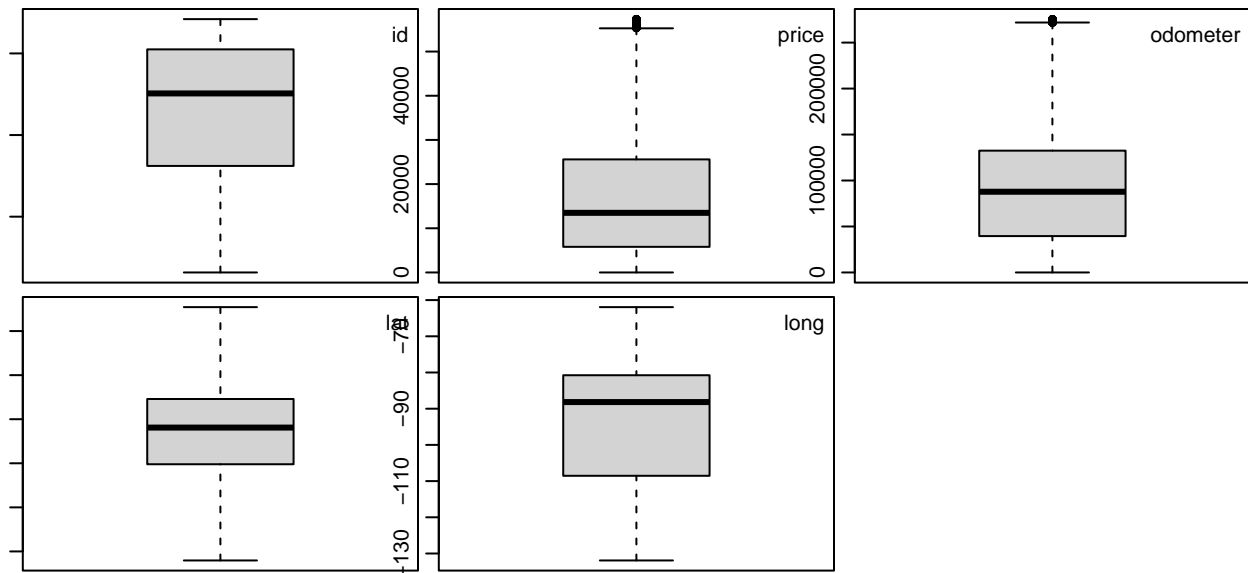
    data <- data[data[[col]] >= lower_bound & data[[col]] <= upper_bound, ]
  }
}
return(data)
}

# Remove outliers across all features
vehicles_no_outliers <- remove_outliers_all(vehicles)

# Plot boxplots again to check for outliers removal
plot_boxplots(vehicles_no_outliers)

```





```
# Correlation plot to know about the feature importance
```

```
# Remove outliers from the "price" column
```

```
vehicles <- remove_outliers_all(vehicles)
```

```
# Now, proceed with the label encoding and correlation plot generation as before
```

```
library(corrplot)
```

```
## corrplot 0.92 loaded
```

```
label_encode_df <- function(dataframe) {
  for (col in names(dataframe)) {
    if (is.factor(dataframe[[col]]) || is.character(dataframe[[col]])) {
      dataframe[[col]] <- as.integer(factor(dataframe[[col]]))
    }
  }
  return(dataframe)
}
```

```
# Apply label encoding
```

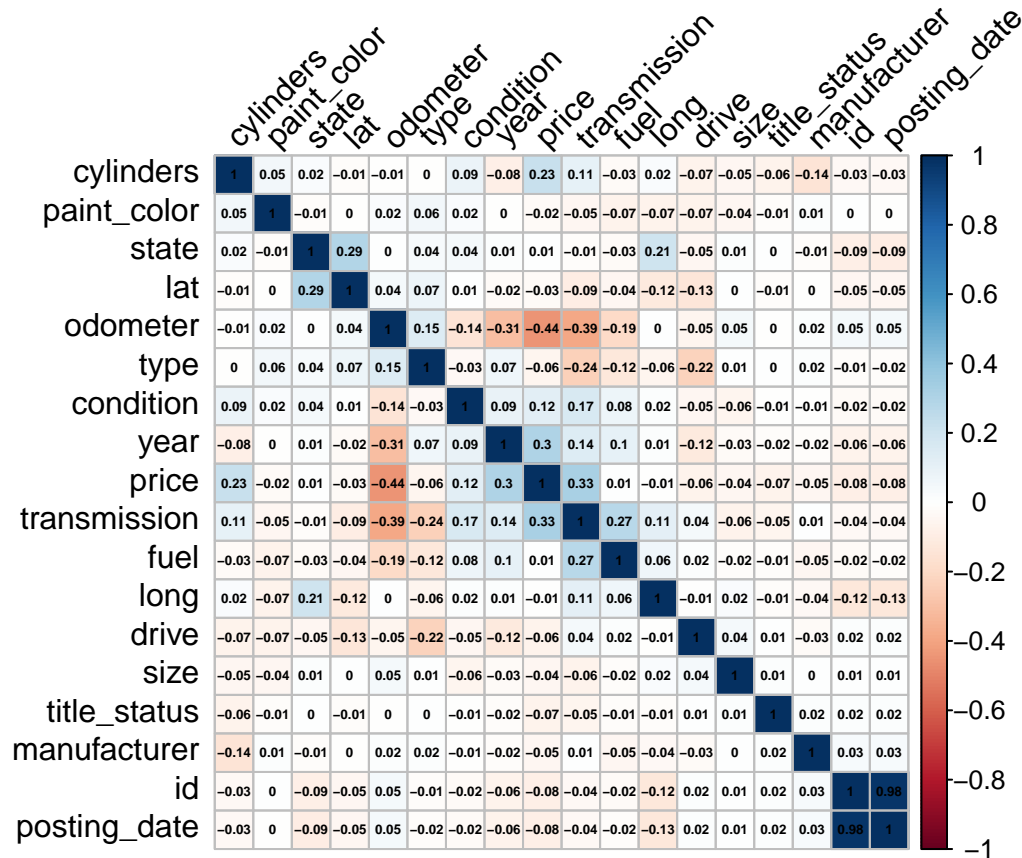
```
df_encoded <- label_encode_df(vehicles)
```

```
# Plot correlation matrix
```

```
par(mfrow = c(1, 1), mar = c(5, 5, 5, 5))
```

```
options(repr.plot.width = 15, repr.plot.height = 15)
```

```
correlation_matrix <- cor(df_encoded)
corrplot(correlation_matrix, method = "color",
  order = "hclust", tl.col = "black", tl.srt = 45,
  addCoef.col = "black", number.cex = 0.45, mar = c(0, 0, 0, 0), addgrid.col = "gray")
```



```
head(vehicles)
```

```
##           id price year manufacturer condition  cylinders fuel odometer
## 28 7316814884 33590 2014          gmc         good 8 cylinders  gas   57923
## 29 7316814758 22590 2010    chevrolet         good 8 cylinders  gas   71229
## 30 7316814989 39590 2020    chevrolet         good 8 cylinders  gas   19160
## 31 7316743432 30990 2017      toyota         good 8 cylinders  gas   41124
## 32 7316356412 15000 2013        ford excellent 6 cylinders  gas  128000
## 33 7316343444 27990 2012          gmc         good 8 cylinders  gas   68696
## title_status transmission drive      size  type paint_color state  lat
## 28      clean          other 4wd full-size pickup      white  al 32.590
## 29      clean          other 4wd full-size pickup      blue  al 32.590
## 30      clean          other 4wd full-size pickup      red  al 32.590
## 31      clean          other 4wd full-size pickup      red  al 32.590
## 32      clean    automatic  rwd full-size truck      black  al 32.592
## 33      clean          other 4wd full-size pickup      black  al 32.590
##           long           posting_date
## 28 -85.4800 2021-05-04T12:31:18-0500
## 29 -85.4800 2021-05-04T12:31:08-0500
## 30 -85.4800 2021-05-04T12:31:25-0500
```



```
## 31 -85.4800 2021-05-04T10:41:31-0500
## 32 -85.5189 2021-05-03T14:02:03-0500
## 33 -85.4800 2021-05-03T13:41:25-0500
```

```
# Step4 : Checking duplicated records
```

```
# Checking duplicated rows
```

```
duplicates <- vehicles[duplicated(vehicles), ]
```

```
# Printing duplicated rows if any
```

```
print(duplicates)
```

```
## [1] id          price          year          manufacturer condition
## [6] cylinders    fuel           odometer       title_status transmission
## [11] drive        size           type           paint_color    state
## [16] lat          long           posting_date
## <0 rows> (or 0-length row.names)
```

The above value of 0 concludes that there are no duplicated records.

```
# Step5 : Transforming "Odometer" feature into categories of (low,medium,high)
```

```
quan25 <- quantile(vehicles$odometer,0.25)
```

```
quan50 <- quantile(vehicles$odometer,0.50)
```

```
a <- function(val){if(val< quan25){
  return('Low')
} else if(val> quan25 & val< quan50){
  return('Medium')
} else{
  return('High')
}}
```

```
vehicles$odometer_status <- sapply(vehicles$odometer,a)
```

```
# Step6 : Transforming "Postingdate" feature by removing the timeframe and just keeping the date as date
```

```
b <- function(val){
  substring(val,1,10)
}
```

```
vehicles$posting_date <- sapply(vehicles$posting_date,b)
```

```
# Step7 : Transforming "cylinders" feature by removing the non-numeric characters and considering only numeric
```

```
vehicles$cylinders <- as.integer(substr(gsub("^\\D*", "", vehicles$cylinders), 1, 1))
```

```
mode <- as.integer(names(sort(table(vehicles$cylinders), decreasing = TRUE))[1])
```

```
vehicles$cylinders[is.na(vehicles$cylinders)] <- mode
```

```
head(vehicles)
```

```
##          id price year manufacturer condition cylinders fuel odometer
## 28 7316814884 33590 2014          gmc          good           8  gas   57923
## 29 7316814758 22590 2010    chevrolet          good           8  gas   71229
```

```
## 30 7316814989 39590 2020 chevrolet good 8 gas 19160
## 31 7316743432 30990 2017 toyota good 8 gas 41124
## 32 7316356412 15000 2013 ford excellent 6 gas 128000
## 33 7316343444 27990 2012 gmc good 8 gas 68696
## title_status transmission drive size type paint_color state lat
## 28 clean other 4wd full-size pickup white al 32.590
## 29 clean other 4wd full-size pickup blue al 32.590
## 30 clean other 4wd full-size pickup red al 32.590
## 31 clean other 4wd full-size pickup red al 32.590
## 32 clean automatic rwd full-size truck black al 32.592
## 33 clean other 4wd full-size pickup black al 32.590
## long posting_date odometer_status
## 28 -85.4800 2021-05-04 Medium
## 29 -85.4800 2021-05-04 Medium
## 30 -85.4800 2021-05-04 Low
## 31 -85.4800 2021-05-04 Medium
## 32 -85.5189 2021-05-03 High
## 33 -85.4800 2021-05-03 Medium
```

```
colSums(is.na(vehicles)==TRUE)
```

```
## id price year manufacturer condition
## 0 0 0 0 0
## cylinders fuel odometer title_status transmission
## 0 0 0 0 0
## drive size type paint_color state
## 0 0 0 0 0
## lat long posting_date odometer_status
## 0 0 0 0
```

```
# Install and load necessary packages
options(repos = "https://cran.r-project.org/")
install.packages("ranger")
```

```
## Installing package into '/Users/narayanaroyal/Library/R/x86_64/4.3/library'
## (as 'lib' is unspecified)

##
## The downloaded binary packages are in
## /var/folders/w8/s5t40zys2lg2hxnvvxxnyqdm0000gn/T//RtmpFwVP6P/downloaded_packages
```

```
library(caret)
```

```
## Loading required package: lattice
```

```
library(ranger)
library(xgboost)
```

```
## Warning: package 'xgboost' was built under R version 4.3.2
```

```

library(rpart)

# Split the data into training and testing sets
index <- createDataPartition(vehicles$price, p = 0.8, list = FALSE)
train_data <- vehicles[index, ]
test_data <- vehicles[-index, ]

# Prepare the data for modeling
x_train <- train_data[, -1] # Exclude 'price' column
y_train <- train_data$price
x_test <- test_data[, -1] # Exclude 'price' column
y_test <- test_data$price

# Convert to data frames
x_train <- as.data.frame(x_train)
x_test <- as.data.frame(x_test)

# Convert factor columns to numeric
convert_to_numeric <- function(df) {
  df[] <- lapply(df, function(x) {
    if (is.factor(x)) as.numeric(as.character(x)) else x
  })
  return(df)
}

x_train <- convert_to_numeric(x_train)
x_test <- convert_to_numeric(x_test)

# Remove non-numeric columns
x_train <- x_train[, sapply(x_train, is.numeric)]
x_test <- x_test[, sapply(x_test, is.numeric)]

# Ensure y_train and y_test are numeric
y_train <- as.numeric(y_train)
y_test <- as.numeric(y_test)

# Remove rows with NA values
train_complete_cases <- complete.cases(x_train) & !is.na(y_train)
test_complete_cases <- complete.cases(x_test) & !is.na(y_test)

x_train <- x_train[train_complete_cases, ]
y_train <- y_train[train_complete_cases]

x_test <- x_test[test_complete_cases, ]
y_test <- y_test[test_complete_cases]

# Remove rows with infinite values
x_train <- x_train[apply(x_train, 1, function(row) all(is.finite(row))), ]
x_test <- x_test[apply(x_test, 1, function(row) all(is.finite(row))), ]

# Combine x_train and y_train for ranger
train_data <- cbind(x_train, price = y_train)

```

```

# Train the random forest model using ranger
model <- ranger(
  formula = price ~ .,
  data = train_data,
  num.trees = 100,
  num.threads = 12, # Use all 12 cores
  importance = 'impurity' # Optional: to calculate variable importance
)
# XGBoost model
xgb_model <- xgboost(
  data = as.matrix(x_train),
  label = y_train,
  nrounds = 100,
  nthread = 12, # Use all 12 cores
  verbose = 0
)
# Decision Tree model
dt_model <- rpart(
  formula = price ~ .,
  data = train_data
)
# Make predictions
# Make predictions
rf_predictions <- predict(model, data = x_test)$predictions
rf_rmse <- sqrt(mean((rf_predictions - y_test)^2))
print(paste("Random Forest RMSE:", rf_rmse))

```

```
## [1] "Random Forest RMSE: 6782.20392081608"
```

```

# Make predictions
xgb_predictions <- predict(xgb_model, as.matrix(x_test))
xgb_rmse <- sqrt(mean((xgb_predictions - y_test)^2))
print(paste("XGBoost RMSE:", xgb_rmse))

```

```
## [1] "XGBoost RMSE: 11.5523166704152"
```

```

# Make predictions
dt_predictions <- predict(dt_model, newdata = x_test)
dt_rmse <- sqrt(mean((dt_predictions - y_test)^2))
print(paste("Decision Tree RMSE:", dt_rmse))

```

```
## [1] "Decision Tree RMSE: 11172.610012344"
```

XGboost model is preferred to use in order to make predictions of the secondhand car price.